

Hongli Zhan

Ph.D. Candidate, UT Austin

Department of Linguistics
The University of Texas at Austin
✉ honglizhan@utexas.edu
🌐 [Website: honglizhan.github.io](https://honglizhan.github.io)
🐙 [Github: github.com/honglizhan](https://github.com/honglizhan)
in [Linkedin: hongli-zhan-761291198](https://www.linkedin.com/in/hongli-zhan-761291198)

Education



Ph.D. in Computational Linguistics

Department of Linguistics, The University of Texas at Austin, Austin, TX

- Advisor: Dr. Junyi Jessy Li
- Committee: Dr. Desmond C. Ong, Dr. David I. Beaver, Dr. Kyle Mahowald

Aug 2021 – Present



B.A. in English Linguistics

School of Foreign Languages, Shanghai Jiao Tong University, Shanghai, China

Sep 2017 – June 2021



Industry Experience



Research Scientist Intern (Part-Time)

MBZUAI Institute of Foundation Models (IFM), Silicon Valley Lab, Sunnyvale, CA

- Hosts: Dr. Rupesh Kumar Srivastava & Varad Pimpalkhute
- Post-training and agents

Jan 2026 – May 2026



Research Scientist Intern

IBM Research, IBM Thomas J. Watson Research Center, Yorktown Heights, NY

- Manager: Dr. Raya Horesh; Mentors: Dr. Muneeza Azmat & Dr. Pin-Yu Chen
- Developed a scalable method for continuous safety governance of LLMs; the pipeline directly supports the safety and trustworthiness of IBM's **Granite** models
- Work in preparation for a first-authored paper & a first-authored U.S. patent

May 2025 – Aug 2025



Research Scientist Intern

IBM Research, IBM Thomas J. Watson Research Center, Yorktown Heights, NY

- Manager: Dr. Raya Horesh; Mentors: Dr. Muneeza Azmat & Dr. Mikhail Yurochkin
- Designed and implemented SPRI, a context-aware, principle-guided framework that enables LLMs to achieve human-level alignment with minimal manual oversight; SPRI forms the core synthetic data generation pipeline for IBM's **Granite Guardian** models
- Work led to a first-authored paper at ICML 2025 [C6] & a first-authored U.S. patent [P1]

May 2024 – Aug 2024

Refereed Publications

* denotes equal contributions

(C stands for conference publications, J for journals, and P for patents)

In Conference Proceedings

- [C6] 2025 **Hongli Zhan**, Muneeza Azmat, Raya Horesh, Junyi Jessy Li, and Mikhail Yurochkin. SPRI: Aligning large language models with context-situated principles. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, 2025. [26.9% acceptance rate (3,260 out of 12,107 submissions)]
- [C5] 2024 **Hongli Zhan**, Allen Zheng, Yoon Kyung Lee, Jina Suh, Junyi Jessy Li, and Desmond C. Ong. Large language models are capable of offering cognitive reappraisal, if guided. In *Proceedings of the 1st Conference on Language Modeling (COLM)*, 2024. [28.8% acceptance rate (299 out of 1,036 submissions)]

- [C4] 2024 Yoon Kyung Lee, Jina Suh, **Hongli Zhan**, Junyi Jessy Li, and Desmond C. Ong. Large language models produce responses perceived to be empathic. In *Proceedings of the 12th IEEE International Conference on Affective Computing and Intelligent Interaction (ACII)*, 2024. [36.7% acceptance rate (44 out of 120 submissions)] 
- [C3] 2023 **Hongli Zhan**, Desmond C. Ong, and Junyi Jessy Li. Evaluating subjective cognitive appraisals of emotions from large language models. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Findings of the Association for Computational Linguistics: EMNLP 2023 (EMNLP Findings)*, pages 14418–14446, Singapore, December 2023. Association for Computational Linguistics. [45.4% acceptance rate (1,758 out of 3,868 submissions)] 
- [C2] 2023 Tiberiu Sosea*, **Hongli Zhan***, Junyi Jessy Li, and Cornelia Caragea. Unsupervised extractive summarization of emotion triggers. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki, editors, *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers) (ACL)*, pages 9550–9569, Toronto, Canada, July 2023. Association for Computational Linguistics. [23.5% acceptance rate (910 out of 3,872 submissions)] 
- [C1] 2022 **Hongli Zhan***, Tiberiu Sosea*, Cornelia Caragea, and Junyi Jessy Li. Why do you feel this way? summarizing triggers of emotions in social media posts. In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang, editors, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 9436–9453, Abu Dhabi, United Arab Emirates, December 2022. Association for Computational Linguistics. [22.1% acceptance rate (715 out of 3,242 submissions)] 

Patents

- [P1] **Hongli Zhan**, Mikhail Yurochkin, Raya Horesh, Adriana Alvarado Garcia, Heloisa Caroline de Souza Pereira Candello, and Muneeza Azmat. Context-aware principle-guided (framework and system) for synthetic data generation. **U.S. Patent Application** [Filed and Pending].

Professional Activities

Organization and Service

- 2023 – 2024 **Organizer**, *Natural Language Learning Reading Group*, UT Austin
Student Volunteer, *ACL 2023*

Reviewing

COLM	Conference on Language Modeling	2025 – 2026
ACL	Annual Meeting of the Association for Computational Linguistics	2025 – 2026
ICRA	IEEE International Conference on Robotics and Automation	2024
EMNLP	Conference on Empirical Methods in Natural Language Processing	2023

Teaching Experience

The University of Texas at Austin

- Fall 2024 **LIN 371: Machine Learning for Text Analysis**, *Guest Lecturer*
Instructor: Prof. Junyi Jessy Li
- Spring 2024 **LIN 371: Machine Learning for Text Analysis**, *Graduate Teaching Assistant*
Instructor: Prof. Junyi Jessy Li
- Spring 2023 **LIN 306: Intro to the Study of Language**, *Guest Lecturer*
Instructor: Seyeon Park
- Fall 2022 **LIN 373N: Machine Learning Toolbox for Text Analysis**, *Graduate Teaching Assistant*
Instructor: Prof. Junyi Jessy Li
- Fall 2021 **LIN 350: Computational Semantics**, *Graduate Teaching Assistant*
Instructor: Prof. Katrin Erk

Mentoring

2023 – 2024 **Allen Zheng**, *B.S. in Computer Science (Turing Honors) & Mathematics, UT Austin*

Fellowships & Awards

IBM Internal Awards

May 2025 **Patent File Award**, *IBM Research*, 867 USD

Awards from The University of Texas at Austin

Fall 2024 **Professional Development Award for Attending COLM 2024**, *UT Austin*, 1,200 USD

Fall 2023 **Professional Development Award for Attending EMNLP 2023**, *UT Austin*, 1,160 USD

Fall 2022 **Professional Development Award for Attending EMNLP 2022**, *UT Austin*, 1,200 USD

Spring 2022 **COLA Supplemental Graduate School Fellowship**, *UT Austin*, 5,000 USD

Awards from Shanghai Jiao Tong University

2021 **Outstanding Graduate**, *Shanghai Jiao Tong University*

2021 **Outstanding Undergraduate Thesis Award**, *Shanghai Jiao Tong University*

Invited Talks

10/11/2023 **Can I Cheer You Up? Psychologically-Grounded Emotional Support: Leveraging LLMs for Targeted Re-Appraisals**, *Computational Affective and Social Cognition Lab (CASCogLab)*, Department of Psychology, The University of Texas at Austin, Austin, TX

02/22/2023 **An Introduction to Computational Linguistics**, *Westlake High School*, Austin, TX